

06/09

A brief history

1. Pre history (1400 - 1930s)

Pascal, Leibniz - calculators

↳ binary system (1705)

Jacquard - weaving machines (end of 18th c.)

Johann Helerich von Müller (1486)

Charles Babbage → Analytical Engine (1837)

↳ concept: arithmetic + memory + I/O + IFs
unit devices conditionals

Ada, countess of Lovelace → first programmer (1840s)

Vannevar Bush (1931) - mechanical integrators

(analog computers)

2. 1st generation: electro-mechanical (1930s - 1960s)

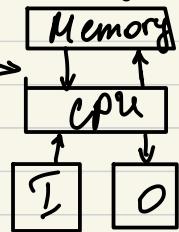
ASCC - Howard Aiken 1937 - 1944

ENIAC - John Mauchly + John Presper Eckert (1943-1945)

first computer SP lawsuit: vs John Atanasoff

EDVAC - Mauchly + Eckert (1948 - 1950)

basis of Von Neumann Architecture



3. 2nd generation: transistors (1965 - 1975)

DEC PDP-1 (1959)

Cray's CDC 6600 - first supercomputer

Sintel 4004 (1971)

Apple (1976)

first monitor (1951)

first mouse (1968)

IBM personal computer - released in 1980s
open standards to third-party

Microprocessors (1960s - today)

4. IV generation: multi-computing
ARPANET (1965-1969) - Leonard Kleinrock
1972 - ARPANET public + Email

08/09

Assembly

Logic circuits [3, 4]

Data representation [5, 6]

Instruction Set Architectures [7, 8]

Software [2]

- assembly

- C

- C++ / Java

13/09

Logic Circuits §

$$x \cdot x = x^2 = x$$

$$x \wedge x = x$$

$$f(x) = a_0 + a_1 x + a_2 x^2 + \dots$$

$$= a_0 + a_{m_2} x$$

$$x + x = 2x = x$$

$$x \vee x = x$$

$$f(x) = a \cdot x + b \cdot (1 - x)$$

$$f(x) = f(1) \cdot x + f(0) \cdot \bar{x}$$

$$1 - x = \bar{x}$$

$$\neg x$$

$$f(x, y) = f(1, 1) \cdot \boxed{x} \cdot \boxed{y} + f(1, 0) \cdot \boxed{x} \cdot \boxed{\bar{y}} +$$

minterms

$$f(0, 1) \cdot \boxed{\bar{x}} \cdot \boxed{y} + f(0, 0) \cdot \boxed{\bar{x}} \cdot \boxed{\bar{y}}$$

XOR

x	y	$x \otimes y$
0	0	0
1	0	1
1	1	0

$$f(x, y) = f(1, 1) \cdot x \cdot y + f(0, 1) \cdot \bar{x} \cdot y + f(1, 0) \cdot x \cdot \bar{y} + f(0, 0) \cdot \bar{x} \cdot \bar{y}$$

0	1	1
1	0	1
1	1	0

$$x \otimes y = f(0, 1) \cdot \bar{x} \cdot y + f(1, 0) \cdot x \cdot \bar{y}$$

$$x \wedge y = f(1, 1) \cdot x \cdot y$$

$$f(x, y, z) = (\bar{x} \cdot \bar{z}) \cdot y + \bar{z} \cdot \bar{x}$$

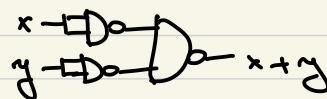
$$= (x + z) \cdot y + \bar{z} \cdot \bar{x} = xy + z \cdot y + \bar{z} \cdot \bar{x}$$

$$\bar{x} \cdot \bar{y} = \bar{x} + \bar{y}$$

$$\bar{x} + y = \bar{x} \cdot \bar{y}$$

$$x \cdot \bar{y} = x$$

x	y	z	f
0	0	0	1
0	0	1	0
0	1	0	1
0	1	1	1
1	0	0	0
1	0	1	0
1	1	0	1
1	1	1	1



$$x \quad y \quad f \quad f = x \cdot y + \bar{x} \cdot y + x \cdot \bar{y}$$

$$0 \quad 0 \quad 0$$

$$0 \quad 1 \quad 1$$

$$1 \quad 0 \quad 1$$

$$1 \quad 1 \quad 0$$

$$f = \bar{x} \cdot \bar{y} \cdot \bar{z} + \bar{x} \cdot \bar{y} \cdot z + \bar{x} \cdot y \cdot \bar{z} + x \cdot y \cdot z$$

$$= (\bar{x} \cdot \bar{y})(\bar{z} + z) + y \cdot z(\bar{x} + x)$$

$$= \bar{x} \cdot \bar{y} + y \cdot z$$

$\bar{z} \cdot \bar{y}$	00	01	11	10
0	1	0	0	0
1	1	1	1	0

$\bar{z} \cdot \bar{y}$	00	01	11	10
00	1	1	0	1
01	1	1	0	1
11	0	0	0	0
10	1	1	1	0

$$f = \bar{x} \cdot \bar{z} + \bar{y} \cdot z + y \cdot z$$

$\bar{z} \cdot \bar{y}$	00	01	11	10
0	1	1	0	0
1	1	1	1	1

15/09

Logic Circuits §

20/09

Data Representation §

$\mathbb{Z}^+ \rightarrow$ nonnegative integers
 $\Sigma = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

$2557_{10} \rightarrow$ binary (radix-2)
2 5 5 7
0010 0101 0101 0111

-52_{10} 8 bit 2's comp $52_{10} \rightarrow 00110100$
 $-64_{10} \rightarrow 1000 0000$ ↓
 $\underline{-64_{10}}$ $\begin{cases} 1000 \\ 1100 \end{cases}$ $-52_{10} \rightarrow 11001100$
 12_{10}
 64_{10}

4998_{10} in BCD

$4998 \div 2 = 2499$ 0 1
 $2499 \div 2 = 1249$ 1 2
 $1249 \div 2 = 624$ 1 4
 $624 \div 2 = 312$ 0 8
 $312 \div 2 = 156$ 0 16
 $156 \div 2 = 78$ 0 32
 $78 \div 2 = 39$ 0 64
 $39 \div 2 = 19$ 1 128
 $19 \div 2 = 9$ 1 256
 $9 \div 2 = 4$ 1 512
 $4 \div 2 = 2$ 0 1024
 $2 \div 2 = 1$ 0 2048
 $1 \div 2 = 0$ 1 4096

22/09

Data Representation

 $S \& M \rightarrow$ overflow not detected

$$\begin{array}{r}
 + 0111 \quad (+7) \leftarrow \text{bigger magnitude} \\
 + 1011 \quad (-3) \\
 \hline
 0100 \quad (+4)
 \end{array}$$

 $1C \rightarrow$ overflow not detected

$$\begin{array}{r}
 + 0111 \quad (+7) \\
 + 1100 \quad (-3) \rightarrow 0011_2 = 3_{10} \\
 \hline
 \begin{matrix} \text{carry out} \\ 1 \end{matrix} \quad 0011 \quad 1+ \\
 0100 \quad (+4)
 \end{array}$$

 $2C \rightarrow$ overflow detected iff MSB of both is the same,

$$\begin{array}{r}
 + 0111 \quad (+7) \quad \text{the sum should have the} \\
 + 1101 \quad (-3) \rightarrow 0011_2 = 3_{10} \quad \text{same MSB (0 \rightarrow + | 1 \rightarrow -)} \\
 \hline
 \begin{matrix} \text{carry out} \\ 1 \end{matrix} \quad 0100
 \end{array}$$

integer values $105949 \div 105989$ \hookrightarrow Excess -105949

4224 for use in LCD Clock

 \hookrightarrow BCD

Fixed point fractions

$$x \cdot 2^{-x} \Rightarrow 8.45_{10}$$

$$0 \frac{1}{2} \downarrow$$

$$1 \quad 0.5 \quad 1000.11_2$$

$$2 \quad 0.25$$

$$3 \quad 0.125$$

 $\dots \dots$

$$e=0, \text{ value}=0 | e \neq 2^{15} \text{ value}=\pm\infty, m=0$$

Floating point numbers

sign	exponent	mantissa
s	e $30 \xleftarrow{8 \text{ bits}} 23 \dots 22 \xleftarrow{23 \text{ bits}} 0$	m $1. m$

$$\text{value} = (-1)^s \times 2^{e-127} \times m$$

$$m' = 1.m$$

$$1 \leq e \leq 254 \quad \text{excess } -127 \quad [-126; 127]$$

$$\begin{aligned} \text{value} &= 8.75 \\ \text{fixed: } &1000.11 \\ &1.00011 \times 2^3 \end{aligned}$$

$$\begin{array}{c|c|c|c} s & e & & m \\ \hline 0 & \underbrace{10000010}_{+130-127=3} & & \underbrace{000110000\dots}_{m' = 1.00011} \end{array}$$

$$\begin{array}{c|c|c|c} s & e & & m \\ \hline 1 & 10000000 & 0101000\dots & = -2.625_{10} \\ \hline \downarrow & 128-127=1 \Rightarrow 2^1 & 1.0101 & \Rightarrow 10.101_2 = 2_{10} + 0.5_{10} + 0.125_{10} = 2.625_{10} \end{array}$$

Sign	Exponent	Mantissa
0 +	10000010 3	001001100000000000000000 $\xrightarrow{+} 1001.00111$
0 +	10000011 4	001010001100000000000000 $\xrightarrow{+} 10010.100011$
0	10000011	1011110001 $\xrightarrow{+} 1.1011.110001$

	s	exponent	mantissa	
+	0	10000011	001101101000...	111111 010011.01101 +
	0	10000100	001011111000...	100101.1111
	0	10000100	11001010110000...	111001.01011

22/09

Instruction Set Architecture 3

stack based \rightarrow VISC

register based \rightarrow RISC

VISC \rightarrow very long instruction word

RISC \rightarrow minimal instruction set computer

Memory Layout

bits \rightarrow flip-flop

words \rightarrow parallel access register

memory chip (SoC)

byte addressable

$KB = KiB$ (binary kilobyte) = 1024 bytes

$$4 \text{ GB} = 2^{10} \cdot 2^{10} \cdot 2^{10} \cdot 2^2 \rightarrow 32 \text{ bytes} \rightarrow 35 \text{ bits} \rightarrow 30 \text{ words of 32 bits}$$

big endian memory: ...|a|b|c|... \rightarrow can be processed as
little endian memory: ...|c|b|a|... \rightarrow waits for last one

32 registers \rightarrow 5 bit reg operand

3-reg instruction \rightarrow 15 bit $\Rightarrow 32 - 15 = 17$ bit OP

reg+dir in 16 MB memory
 $\hookrightarrow 2^9$ $\downarrow 2^{24}$

$\hookrightarrow 32 - 2^9 = 3$ bit OP $\Rightarrow 2^3 = 8$ instructions

6⁴ reg + reg + imm + ? 16 bit const
 $\hookrightarrow 2^6$ $\downarrow 2^6$ $\downarrow 2^5$ $\downarrow 2^5$ $\downarrow 16 = 32$ ✓

Structure Set Architecture

register transfer notation

[R1] - register

M(A) - memory (direct)

TOS - top of stack

29/09

Accumulator = Memory + TOS (Processor)

Stack = TOS + Stack (Processor)

Register-Memory = Memory + registers (Processor)

Register-Register = registers (Processor)

1. ACCU-based ISA

single-register

simple design \rightarrow easy to build and program

hard access to data (bottleneck memory)

example: 16-bit words, byte-addressable

S & M integers

instruction \rightarrow 4 OP, 12 operand \Rightarrow 4 KiB

JS: Data copy (load, store), Arithmetic & Logic, control

LOAD 3 take M(3)

SUB 3 sub M(3) M(3)=0

STORE 3 store in M(3)

2. Register-based ISA

general purpose registers

- faster than memory

- fewer address bits

bottleneck memory

EA = effective address

index operand: $EA = [Reg] + I_{dp}$

Scale s = 1, 2, 4 or 8

Displacement disp = 8, 16, 32 or 64 bit

3. Multiple-register Architecture

CISC style: 15 bytes at most

Legacy prefixes (1-4 bytes, optional)

Opcode with prefixes (1-4 bytes)

ModR/M (1 byte, if required)

SIB (1 byte, id required)

Displacement (1/2/4/8 bytes, id required)

Immediate (1/2/4/8 bytes, id required)

RISC (Reduced Instruction Set)

Pro-CISC

easier to program

reduced code size

complex hardware

+ legacy

Con-CISC

complex encoding

may be slow

large energy consumption

KISS (Keep it simple and stupid)

trade-off: CISC vs. RISC

HW vs. SW

11/10

Basic Processing Unit (CPU)

RISC

1. Fetch instr from mem
2. Decode instr and read reg
3. Execute ALU op
4. Memory access
5. Writeback to reg



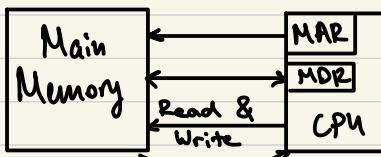
CISC

MAR - Memory Address Register (sends to memory, return)

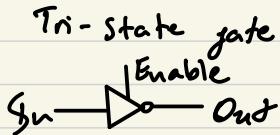
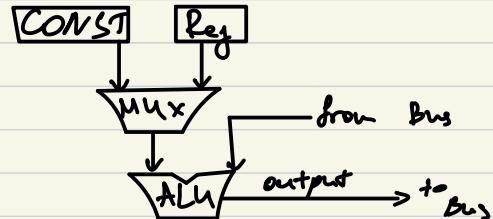
MDR - Memory Data Register (send and receive from mem)

Instruction address generator (PC)

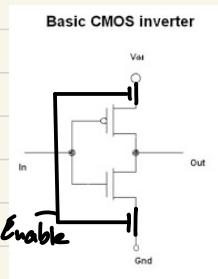
2. Memory
3. MAR / MDR
4. IR (to Control circuitry)



MFC - memory function complete



Enable	In	Out
0	X	Z
1	0	1
1	1	0



ADD R2, (Z1)

Even

1. R₁ out, X_{IN}
2. R₁ out, MAR_{IN}, READ, WMFC (WMFC?)
3. MDR_{out}, ADD, Z_{IN}
4. ~~R₁ out, MAR_{IN}~~
5. Z_{out}, MDR_{IN}, WRITE, WMFC

Wait for Memory
to Complete

1. R₁ out, MAR_{IN}, READ
2. R₂ out, X_{IN}, WMFC
3. MDR_{out}, ADD, Z_{IN}
4. Z_{out}, MDR_{IN}, WRITE, WMFC

13/10

Tutorial

Exam Question 9 - Fixed w/ Floating Point

5 -----. ----- = 11 bits Fixed point

Largest positive - 1.0₁₁₀
 smallest positive - 0.101_{0.101}
 31.5 32 - 0.5 11111.1 → 5 bit mantissa
 $2^{\frac{7}{-2}} \cdot 2^{\frac{-3}{-4}}$ mantissa
 $2^{\frac{7}{-2}} + 6$ exp 1 sign bit
 $2^{\frac{7}{-2}} + 4$ 4 bit exponent
10 bits Floating point

Exam Question 11 - 16 fixing ISA

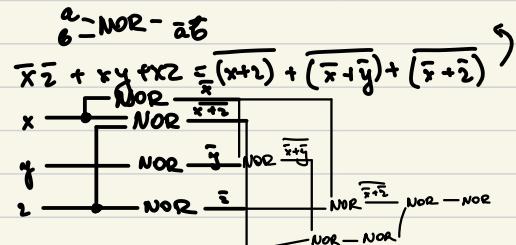
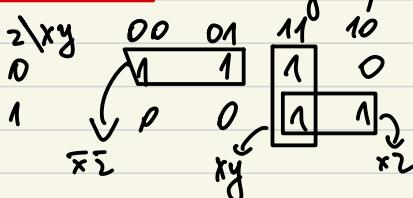
$$16 \text{ MiB} = 2^4 \text{ MiB} = 2^4 \times 2^{20} \text{ B} = 2^{24} \text{ B} \rightarrow 24 \text{ bits addressing}$$

solution - 16 B memory = 2⁴ B → 4-bit dir16 instructions = 2⁴ → 4-bit OP16 registers = 2⁴ → 4-bit reg

mov → ----- 4-bit ----- 4-bit source ----- 4-bit destination
 OP

Question 8 - Bees

$$65536 = 2^{16} \rightarrow \text{overflow of 16 bits}$$

Question 6 - de Morgan / NOR

bits for microcontrolling

2 (reg. add. + 1) → reg.

ex. 16 reg. → 4 reg. add. → $2(4+1) = 10$

ALU → # of op.

ex. 104 → 4 bits

P - only JN-1 bit

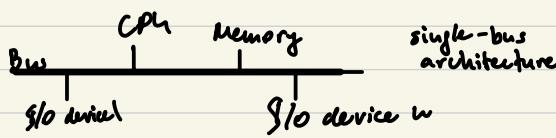
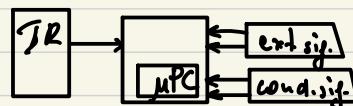
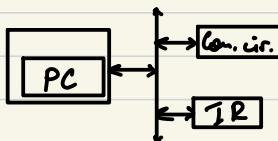
2 - \$N and one - 2 bits

MA2 - only JN-1 bit

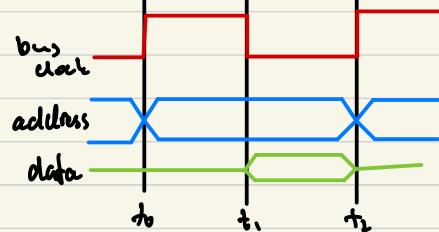
MDR - \$N and one - 2 bits

18/10

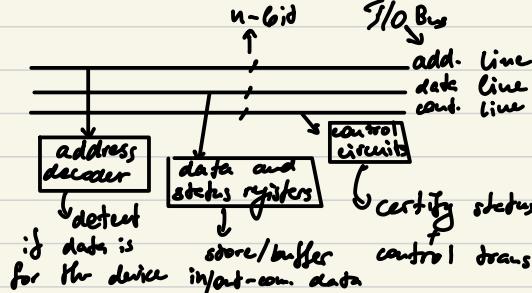
Input/Output



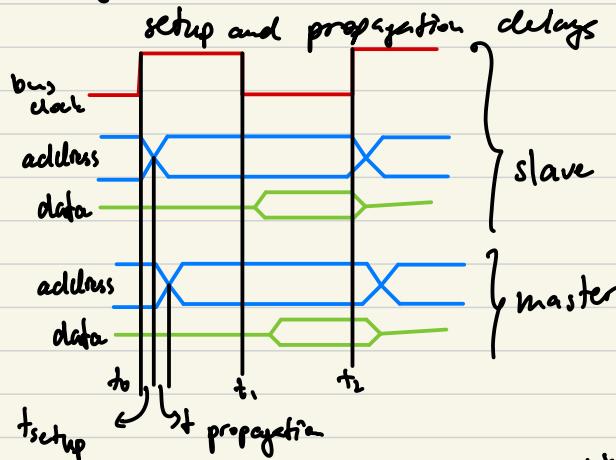
synchronous bus



dual-bus architecture



synchronous bus

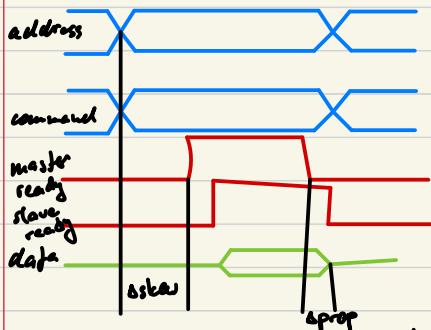


bus driver delay 1.5ns
max bus length 15cm

$$\text{speed of light} \\ c \text{ in copper}$$

$$\frac{c}{2 \times 10^8 \frac{\text{m}}{\text{s}}} = 1.5 \times 10^{-11} \text{ s} = 1.5 \text{ ns}$$

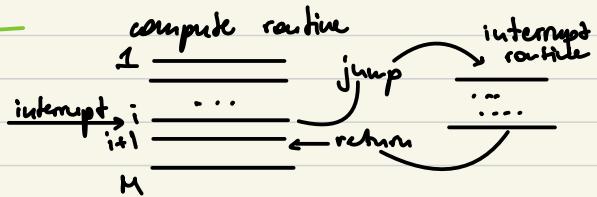
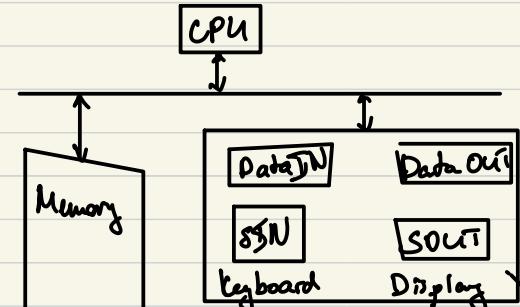
address decoder delay 6ns
time to fetch data $3-15\text{ns}$
setup time (control circuitry) $\leq 2\text{ns}$

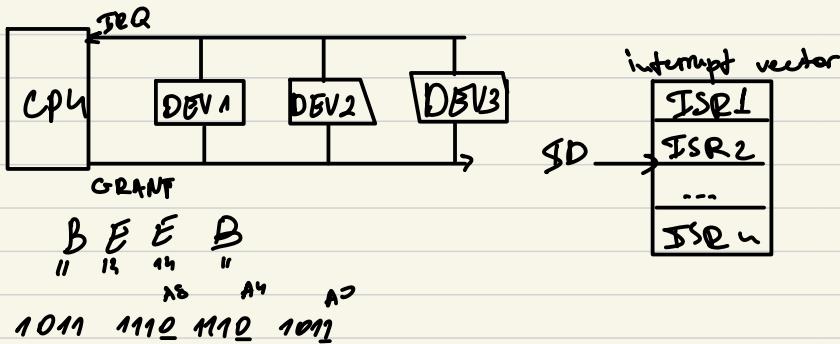


unconditional I/O

passive signaling (polling)

active signaling (interrupt)





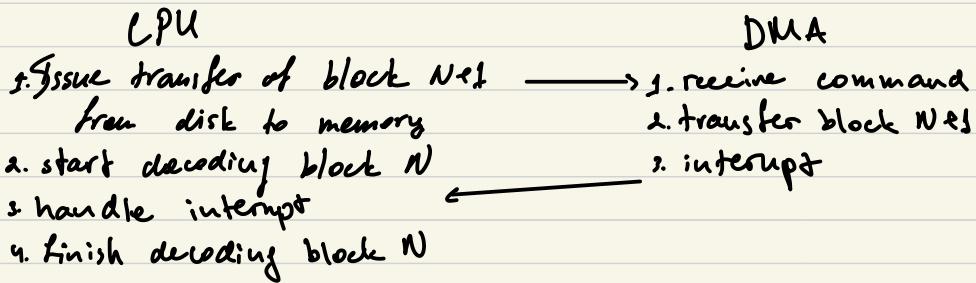
Memory-mapped I/O means I/O devices and memory use the same address space.

20/40

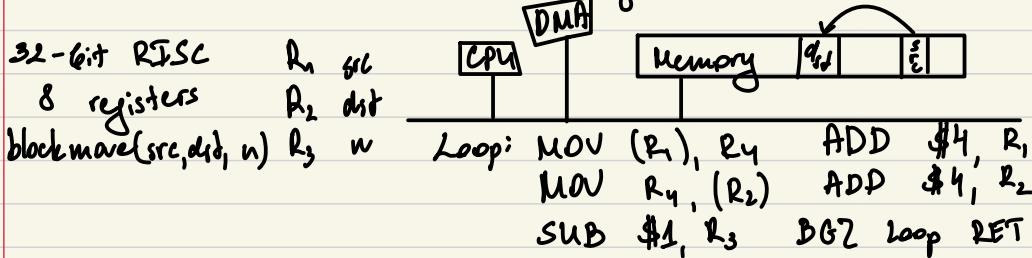
DMA & Memory Organization

DMA - direct memory access

1 instruction: MDV_B (=COPY)



DMA - bus master when transferring data

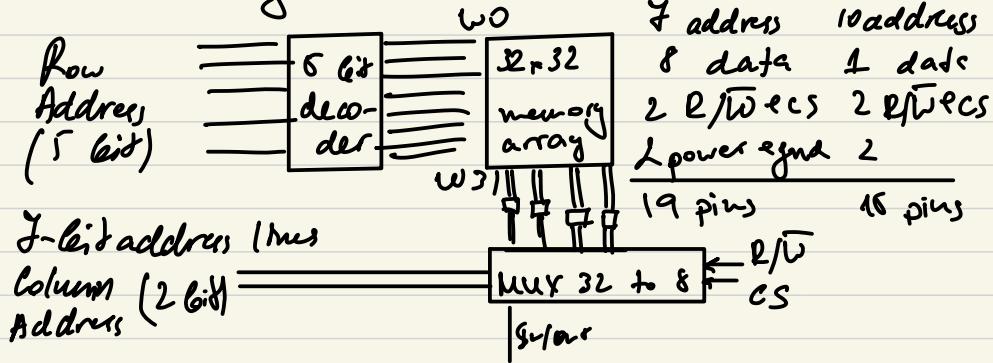


$$\text{speed up} = \frac{\text{time (CPU)}}{\text{time (DMA)}} = \frac{6 \text{ c/w}}{2 \text{ c/w}} = 3 \rightarrow \text{otherwise } 2 \text{ cycle per word}$$

$$= \frac{6 \text{ c/w}}{3 \text{ c/w}} = 2 \rightarrow \text{bus occupied every 3 cycles}$$

$\Rightarrow 3 \text{ cycle per word for DMA}$

1024-6-bit memory - 128×8



32×16
5 address pins
16 data pins
2 R/W + CS
2 power + gnd
25 pins

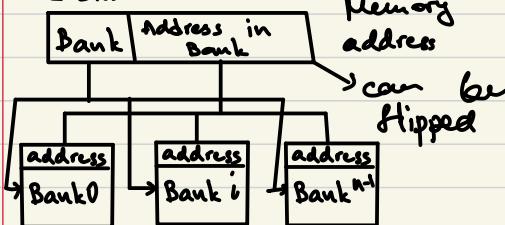
$$64 = 16 \times ? [32]$$

One word = 32 bits

$$MDR = 32^5$$

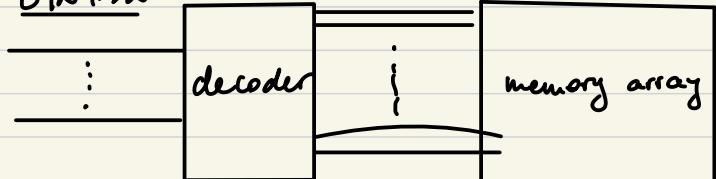
$$MAR = \log_2 32 + 1 = 6$$

? x 32 memory organization
k bits in bits

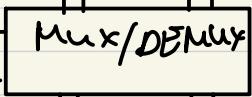


20/10

Tutorial

 512×32 4 row address
lines

$$2048 \text{ bytes} = 2^{11} \text{ bytes}$$

5 column address
line

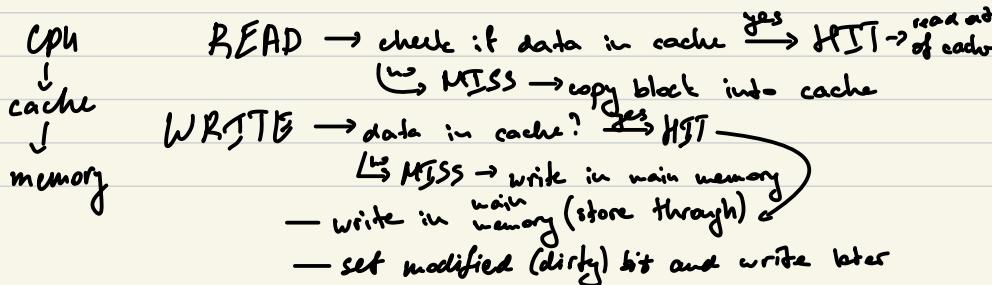
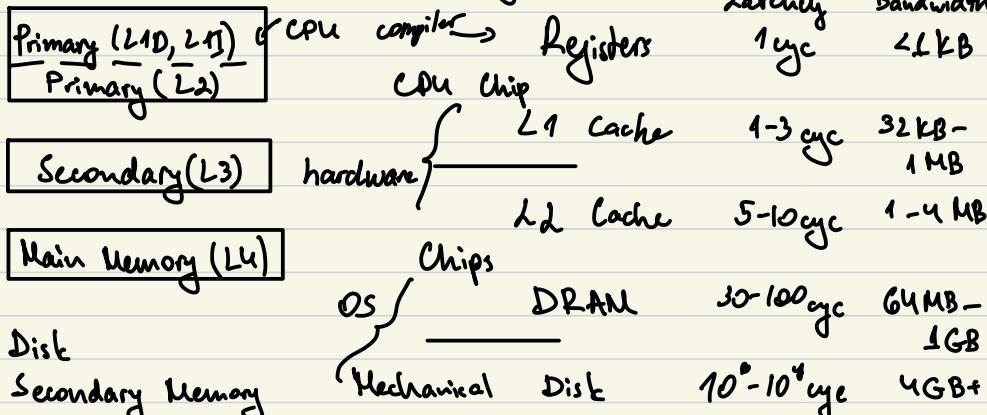
$$\frac{2^4}{2^1 \text{ data}} = 2^3 \text{ address pin} \rightarrow 9 - 4 = 5 \text{ row column}$$

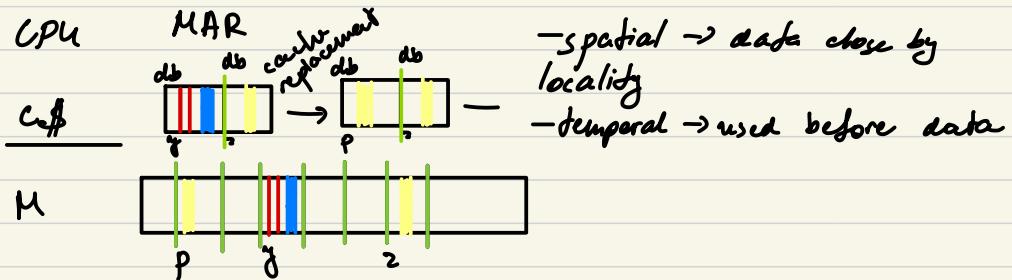
$$2^5 = 32 \text{ data pins}$$

double word

25/10

Caching





L1 cache

$$\text{cache hit ratio: } 0 \leq h \leq 1 \quad \text{avg.} = hc + (1-h)(cm)$$

$$\text{access time ratio: } c$$

$$\text{cache miss ratio: } 1-h \quad \text{speedup} = \frac{\text{time old}}{\text{time new}}$$

$$\text{access time main memory: } m$$

$$h = 0.85 \quad s = \frac{m}{c + (1-h)m} = \frac{400}{40 + 0.15 \cdot 400} = 4$$

$$m = 400 \text{ ns}$$

$$c = 40 \text{ ns}$$

$$m = 400 \text{ ns}$$

$$h = 0.925 \quad s = \frac{m'}{c + (1-h)m} = \frac{800}{40 + 0.045 \cdot 400} = 8$$

$$c = 40 \text{ ns}$$

$$m' = 800 \text{ ns}$$

Cache Mapping

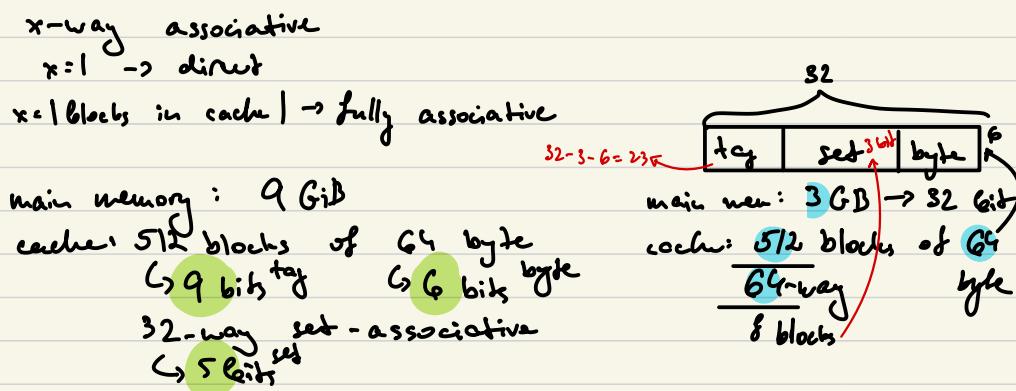
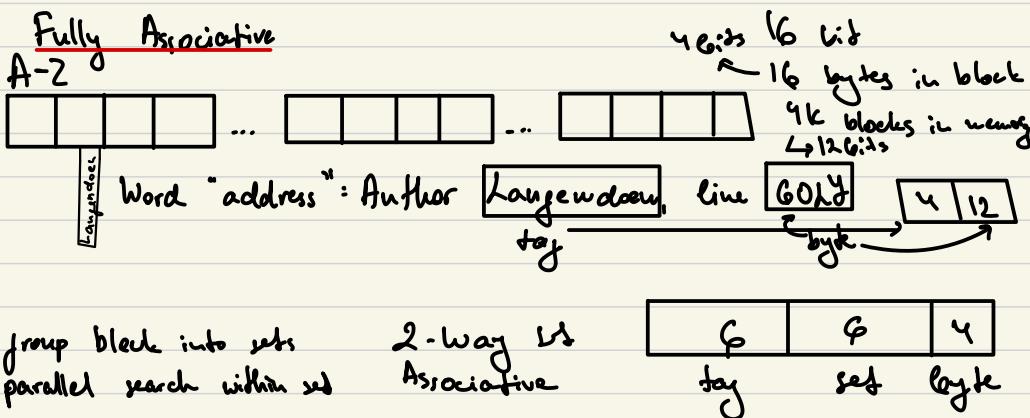
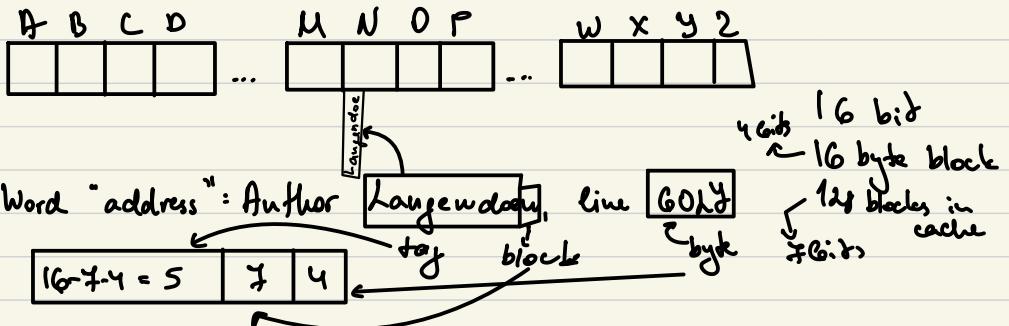
cache - smaller than memory (density)

main mem \rightarrow library

cache \rightarrow books at home (2 books at a time)

block \rightarrow book

byte \rightarrow line



Cash Replacement

x -way associative \rightarrow x -bit age counter per block

HIT

0	1	+1	1	0
0	0	+4	0	1
1	0	=0	0	0
1	1	+1	1	1

age before after

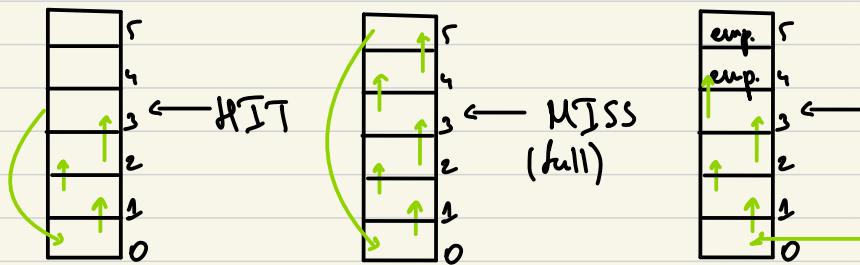
increased - Reset hit counter
 increased of hit block
 w.r.t top - Increment lower
 unchanged counters

MISS

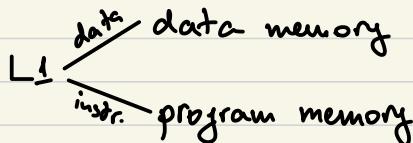
0	1	+1	1	0
0	0	+1	0	1
1	0	+1	1	1
1	1	=0	0	0

age before after

full:
 increased - Replace block with
 increased highest counter, reset
 increased - Increment all
 w.r.t top others
 not full: - empty spot with 0 count
 - Increment all others



Harvard Architecture:
 instructions - spatial + temporal
 data - temporal + spatial



27/10

Pipelining

- 1 washer \rightarrow 30 min
1 drier \rightarrow 50 min
1 folding mat \rightarrow 40 min
4 people

$$30 \quad 50 \quad 40 = 120 \text{ (one load)}$$

$$30 \quad 50 \quad 40$$

$$30 \quad 50 \quad 40$$

$$30 \quad 50 \quad 40$$

$$\underline{30 + 4 \cdot 50 + 40 = 30 \quad 50 \quad 50 \quad 50 \quad 50 \quad 40 = 270}$$

$$\text{wash} = 30 \text{ min}$$

$$30 \quad 40 \quad 20 = 90 \text{ (one load)}$$

$$\text{dry} = 40 \text{ min}$$

$$30 \quad 40 \quad 20$$

$$\text{fold} = 20 \text{ min}$$

$$30 \quad 40 \quad 20$$

$$\underline{30 + 4 \cdot 40 + 20 = 210 = 30 \quad 40 \quad 40 \quad 40 \quad 40 \quad 20}$$

Performance

Latency (L)

- time for one instruction to finish
- lower is better

Throughput (T)

- the # instructions per time unit
- higher is better
- time per instruction decreases (on average)

Increase clock frequency \rightarrow latency boost

problem: power wall

Multiple threads & cores \rightarrow throughput boost

problem: parallel programming

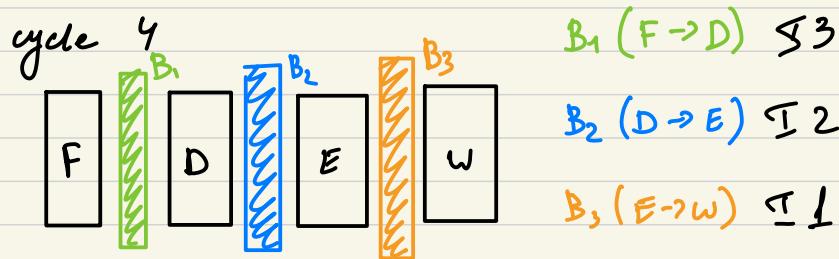
Fetch	Fetch 1	Fetch 2	Fetch 3	Fetch 4
CPU				
Execute time →	Execute 1	Execute 2	Execute 3	

Fetch instruction

Decode instruction and fetch operands

Execute instruction

Write result



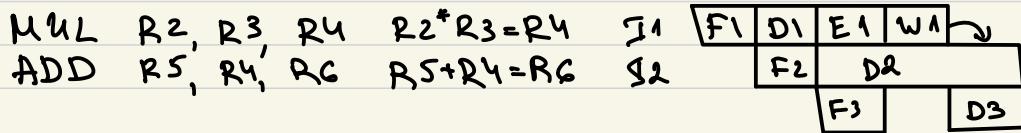
stalls:

Cache miss [fetch, decode]

Dependency between instructions [decode]

Branching [exec]

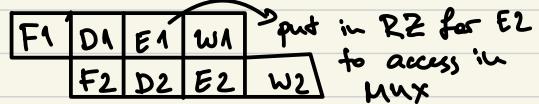
Long operation [exec]



ADD R1, R2 R2 += R1
 MUL R3, R4 R4 *= R3
 ADD R2, R3 R3 *= R2

MOU R1, R2 R2 = R1
 ADD \$100, R2 R2 += 100

ADD R1, R2 R2 += R1
 ADD R2, R3 R3 += R2



BG2 LOOP
 MOV R3, SUM



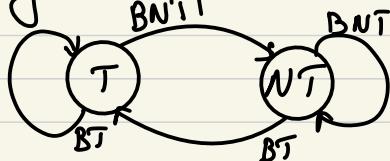
ADD R6, RF, R8
 BG2 R3 LOOP
 MOV R8, SUM

→ BG2 R3 LOOP
 ADD R6, RF, R8
 MOV R8, SUM

Static prediction

backwards branches likely to be taken (loop)
 forward branches mostly skipped (rule vs. exception)
 sign bit (or compiler directed)

Dynamic prediction



ADD R2, R3
DIV R3, R1

I¹
I²
I³
I⁴

F ¹	D ¹	E ¹	W ¹	ADD
F ²	D ²	E ²	W ²	DIV
F ³	D ³	E ³	W ³	
F ⁴	D ⁴	E ⁴	W ⁴	

Execution time of a program: T

Instruction count: N

#cycles per instruction, CPS: $S(y_1, y_2)$

Clock rate: $R(y_1, 1 \text{ GHz})$

without pipelining: $T = (N \times S) / R$

with n-stage pipeline: $T' = T / n$

but S increases due to stalls

$$h_i = 0.95 \rightarrow \text{miss} = 0.05$$

$$h_d = 0.9 \rightarrow \text{miss} = 0.1$$

cache miss $\rightarrow 32$ cycles

20% memory access = 0.2

$$0.35 \times 32$$

$$m_i = 0.05$$

$$m_d = 0.1$$

$$f_{\text{cache}} = 32$$

$$f_{\text{data}} = 0.2$$

$$h_b = 0.8 \rightarrow \text{miss} = 0.2$$

1 additional cycle

20% Branching $\rightarrow 0.2$

$$0.2 \times 0.2 \times 1$$

$$m_{\text{branch}} = 0.2$$

$$f_{\text{branch}} = 1$$

$$d_{\text{branch}} = 0.2$$

Tutorial

$$h_{\text{branch}} = 0.90 \quad m_{\text{branch}} = 0.10$$

$$p_{\text{branch}} = 1 \text{ cycle}$$

$$f_{\text{branch}} = 0.20$$

$$\begin{aligned} \text{Cache misses} &= (m_{\text{instr.}} + m_{\text{data}} \times m_{\text{data}}) \times f_{\text{cache}} \\ &= 1.28 \end{aligned}$$

Branching

$$\begin{aligned} &= p_{\text{branch}} \times m_{\text{branch}} \times p_{\text{branch}} \\ &= 0.2 \times 0.1 \times 1 = 0.02 \end{aligned}$$

01/11

Superscalar Execution

J₁ Fadd

F1	D1	E1	W1	
F1	D2	E2	W1	→ W1

J₃ FsubJ₄ Sub

F1	D1	E3	W1	W1
F1	D2	E2	W1	→ W1

Amdahl's Law

Parallelization can:

- increase throughput
- latency remains same
- reduce program runtime T

not everything can be parallelized

→ data dependencies

→ sequential algorithms

$f_p \rightarrow$ fraction of code that can be parallelized

$1 - f_p = f_s \rightarrow$ sequential execution

$T_s \rightarrow$ time to run program sequentially

$$T_p = T_s \times (f_s + \frac{f_p}{P})$$

$P \rightarrow$ number of processors

$$S = \frac{1}{f_s + \frac{f_p}{P}}$$

$$T_s = 10s$$

$$f_p = 80\% \quad f_s = 20\%$$

$$P = 4$$

$$T_p = 10 \left(20\% + \frac{80\%}{4} \right) = 10.40\% = 4$$

$$S = \frac{T_s}{T_p} = \frac{10}{4} = 2.5$$

1 core with 1 thread

$$f_p = 20\% \quad f_s = 80\%$$

$$P = 8$$

$$S_{max} = \frac{1}{f_s} = \frac{1}{80\%} = 1.25$$

$$S = \frac{T_s}{T_p} = \frac{T_s}{T_s(1 + \frac{f_p}{P})} = \frac{1}{80\% + \frac{20\%}{8}} = \frac{1}{82.5\%} = \frac{100}{82.5} = \frac{20}{16.5}$$

40

20

33

$$s' = \frac{10}{q} = \frac{1}{\frac{fs + fp}{p}} = \frac{1}{\frac{0.8 + 0.2}{p}} = \frac{p}{1.0} = p$$

$$p = ?$$

$$0.8 + \frac{0.2}{2} = 0.9$$

A certain program can be parallelized for 20%
 $fp = 20\% = 0.2$ $fs = 80\% = 0.8$

What is the speedup if we run the program on 8 processors? $p = 8$ $s = \frac{1}{\frac{fs + fp}{p}} = \frac{1}{\frac{0.8 + 0.2}{8}} = \frac{40}{33}$

What is the maximum speedup that can be obtained? $p = \infty$ $s_{max} = \frac{1}{fs} = 1.25$

How many processors are required to get a speedup of $\frac{10}{9}$? $p = ?$ $s' = \frac{1}{\frac{0.8 + 0.2}{p}} = \frac{10}{9}$ $0.9 = 0.8 + \frac{0.2}{p}$ $p = 2$

$$fp = 90\% \quad fs = 10\%$$

$$p = 3 \quad s = \frac{1}{\frac{fs + fp}{p}} = \frac{1}{\frac{10\% + 90\%}{3}} = \frac{3}{10} = 0.3$$

$$s_{max} = \frac{1}{fs} = \frac{1}{fs} = 10 \quad p = ? \quad s = 5 = \frac{1}{\frac{10\% + 90\%}{p}} \quad 0.2 = 0.1 + \frac{0.9}{p}$$

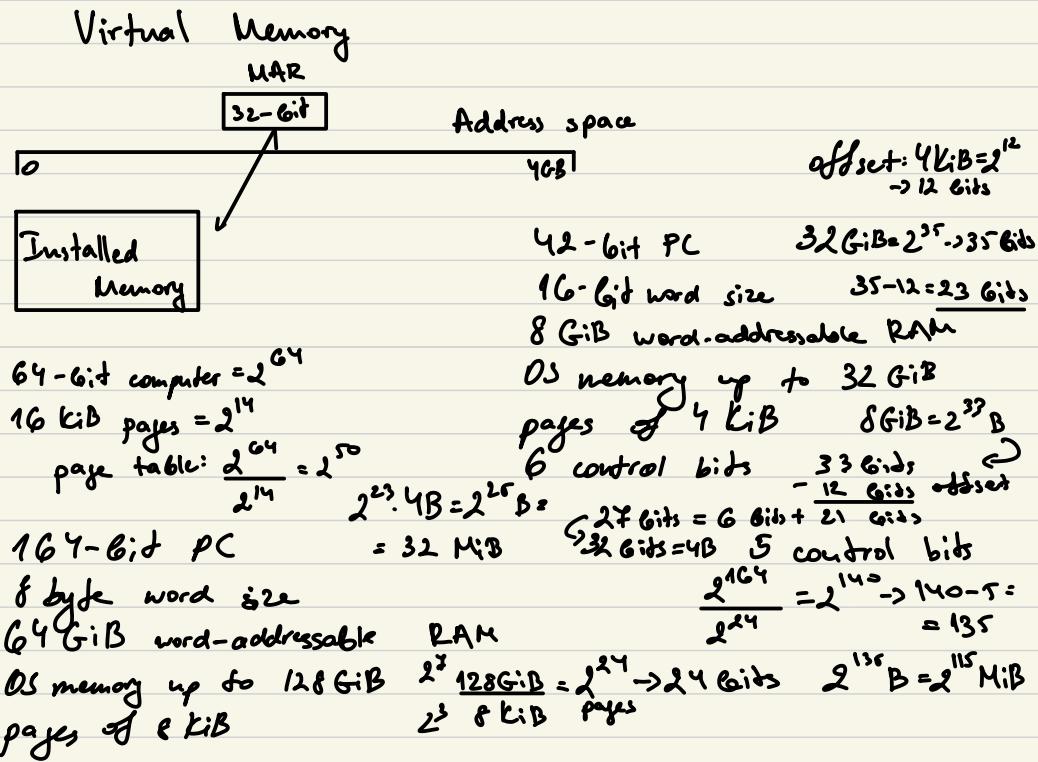
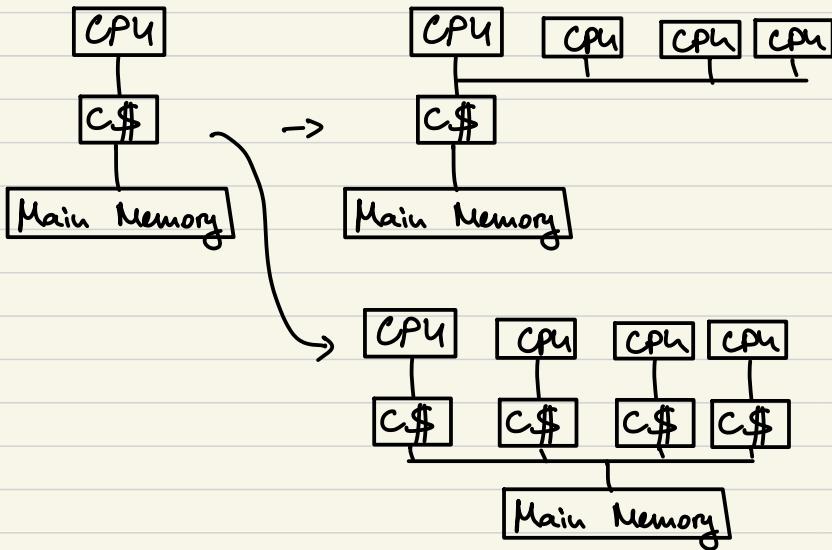
Flynn's Taxonomy

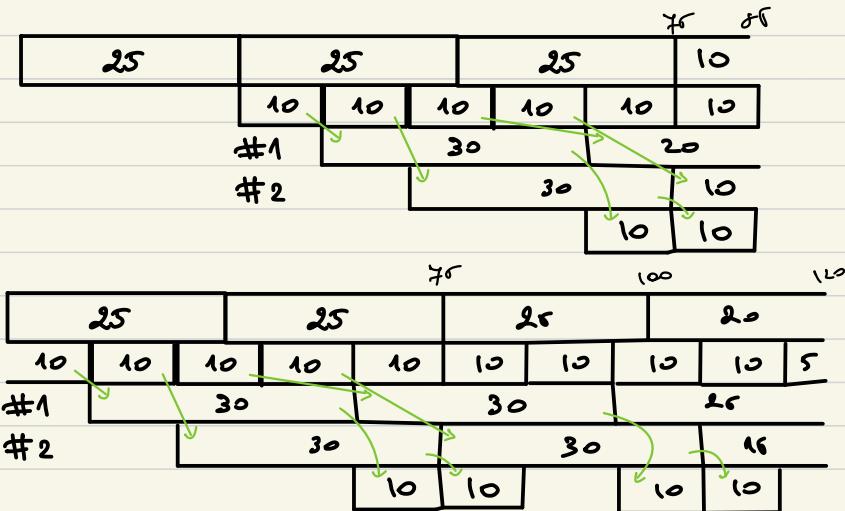
Single Instruction, Single Data (SISD)
 → conventional system

Single Instruction, Multiple Data (SIMD)
 → one instruction on multiple data streams

Multiple Instruction, Multiple Data (MIMD)
 → multiple instruction streams on multiple data streams

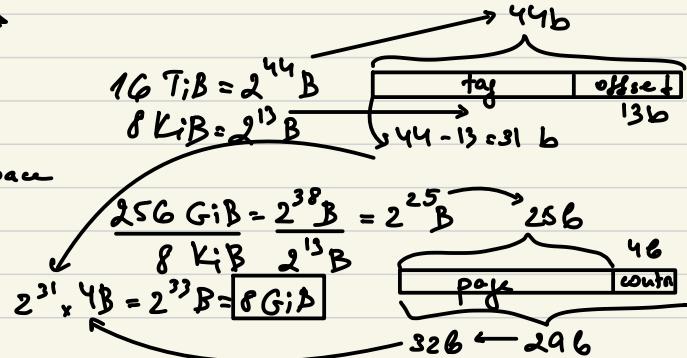
Multiple Instruction, Single Data (MISD)





Virtual Memory

64-bit processor
256 GiB memory
16 TiB address space
8 KiB pages
byte-aligned
4 control bits
page table size = ?



DPU

X, Bus → ALU inputs

memory access ≥ 1 cycle

R2 → register addressing

(R1) → indirect register addressing

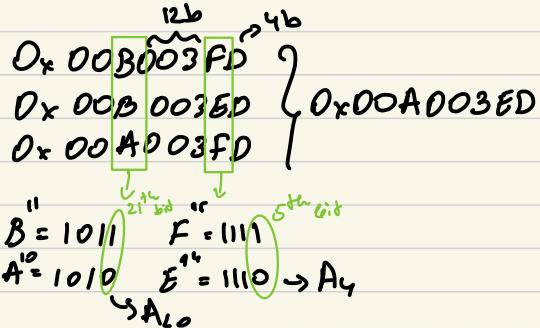
$$a) R1 \leftarrow R2 + (R1)$$

1. R1out, MARin, READ
2. R2out, Yin, WMFC
3. MDRout, Xout, ALUadd, Zin
4. Zout, R1in, END

b) least amount of time

WMFC happens after last input reading

32 address lines
 3 devices connected
 4 ports
 2 ignored address lines



Block 0 hit I_{y}
 Block 3 hit II_{y}
 Block 5 hit I_{y}
 cache miss (full) IV_{y}

instructions

block	before	I _y	II _y	IV _y	after
0	001	010	000	001	010 \rightarrow hit
1	010	011	011	100	101
2	101	101	101	110	111
3	110	110	110	000	001 \rightarrow hit
7	100	100	100	101	110
5	011	000	001	010	011 \rightarrow hit
6	000	001	010	011	100
7	111	111	111	111	000 \rightarrow cache miss

