

Graph Processing

graph - hierarchical data structure modelling dependency relationship

Representation

graph (G) comprises nodes (V) and edges (E)

Adjacency matrix: $n \times n$ matrix $M(n=|V|)$, non-zero M_{ij} means edge from V_i to V_j

Adjacency list: $List[(V, List[V])]$, tuple represents node and its connections

Edge list: $List[(V, V)]$, node pairs representing edge

Structures

Graph components - subgraphs, any two vertices connected by path

Strongly connected comp. - largest sub-graph with path between any two nodes

Triangles/Polygons - three vertices connected to one another

Spanning tree - sub-graph of all nodes and minimum number of edges

Algorithms

Traversal - starting from node, find all connected nodes

Depth-first - recursively follow edges until visiting all reachable nodes

Breadth-first - follow graph edges per level

Node importance - calculate importance of node relative to other nodes

Centrality measures or PageRank

Shortest paths


Dijkstra's algorithm or 'travelling salesman'

Graph DB - specialised RDBMs storing recursive data structures, supporting

CRUD operations on them and maintaining transactional consistency

BSP (bulk synchronous parallel) model

- multiple processors with fast local memory
- pair-wise communication between processors
- barrier implementation synchronising super steps




BSP superstep

- local execution: use own memory to compute on local data partitions
- data exchange/remote communication: data between processes
- barrier synchronisation: processes wait until all finished communicating

Termination

- Superstep 0: all vertices are active
- All active vertices participate in computation at superstep
- Vertex deactivates by voting to halt
- No execution in subsequent supersteps
- Vertex reactivated by receiving a message



Pregel (Google)

- distributed graph processing framework
- computations = sequence of supersteps
- superstep - framework invokes user-defined function for each vertex
- function specifies behaviour at a single vertex (V) and a single superstep (S)
 - read message sent to V in superstep (S-1)
 - send messages to other vertices, to be read at superstep (S+1)
 - modify state of V and outgoing edges

Vertex centric approach

- Reminiscent of MapReduce
 - user focus on local action
 - each vertex processed independently
- well suited for distributed implementation
 - communication from S to S+1
 - no defined execution order within superstep
 - free of deadlocks and data races

Roles

graphs - adjacency lists, partitioned and distributed using network file system

- leader - mapping between data partitions and cluster node; barrier
- worker - maintains for each vertex
 - adjacency list
 - current calculation value
 - queue of incoming messages
 - state

Superstep

- Workers combine incoming messages for all vertices
 - combinator function updates vertex state
- Termination condition met, vertex vote to exclude oneself
- Vertex updates global aggregator (optional)
- Message passing:
 - receiving vertex local: update message queue
 - else: wrap messages per receiving node, send in bulk

Fault tolerance

- reminiscent model of Spark
- periodically leader instruct worker to save in-memory state to storage
- leader - keep-alive messages to workers; failure detection
- failure - leader reassigns graph partitions to live workers

