

## Lecture 1 - Introduction

Moore's Law:  $2 \times \# \text{transistors}/\text{chip}$  every 1.5 years  
 $\Rightarrow 2 \times \# \text{cores}/\text{chip}$  every 2 years  
 $\Rightarrow$  clock speed - not increasing (possibly decreasing)

### Challenges:

- Hardware - processor technology + interconnection network
- Software - parallel compilers + grid computing environment
- Methods & Algorithms

## Lecture 2 - Parallel Computer Architecture

### Processor / Node Architecture

Von Neumann - program and data in same main memory

#### CPU - Central Processing Unit

- fetch, decode, execute program instructions
- fetch data
- store results

#### Graphics Controller

- render display images

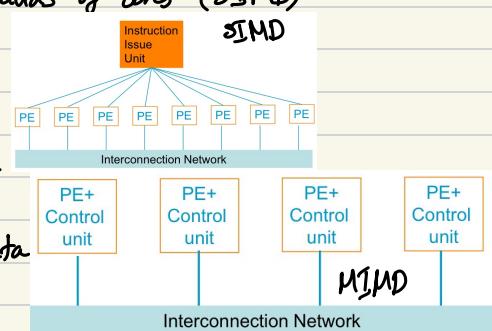
#### GPU - Graphics Processing Unit

- perform complex graphics operations

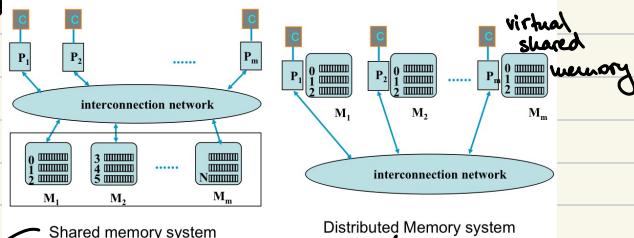
CPU - several cores vs. GPU - thousands of cores (SIMD)

### Flynn's Taxonomy

- 1) Single Instruction, Single Data
- 2) Multiple Instruction, Single Data
- 3) Single Instruction, Multiple Data
- 4) Multiple Instruction, Multiple Data

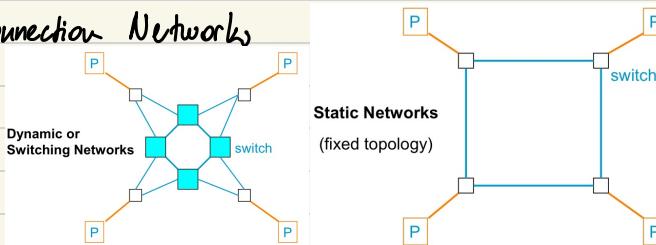


## Memory Organization

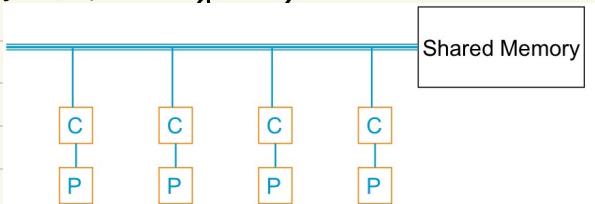


UMA - Uniform Memory Access      NUMA - Non-uniform Memory Access

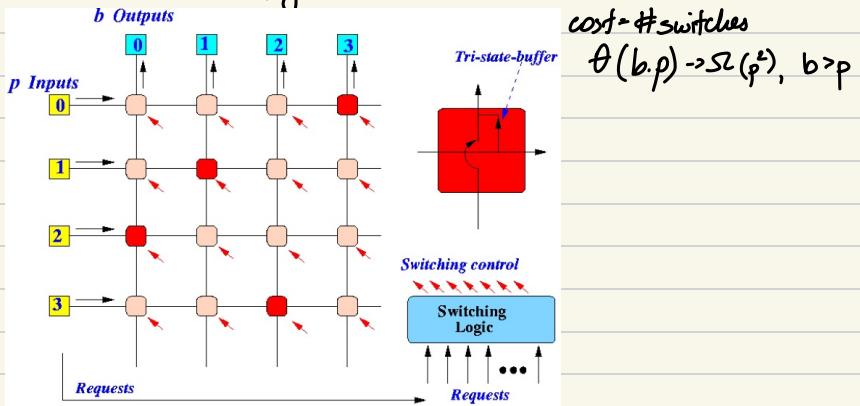
## Interconnection Networks



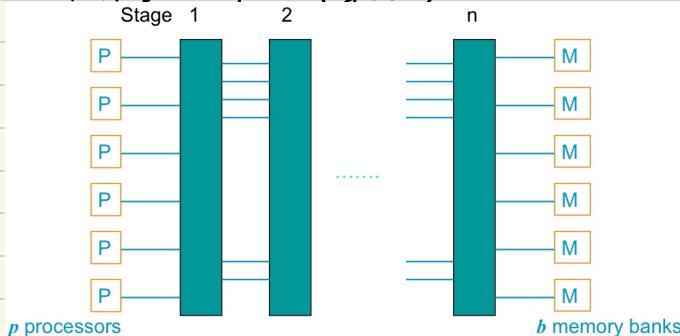
## Bus Network (dynamic)



## Crossbar Network (dynamic)



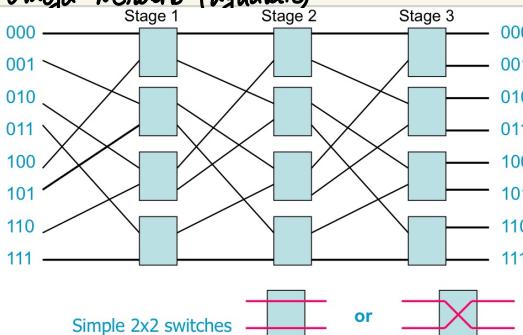
## Multi-stage Network (dynamic)



$$\begin{aligned} \# \text{ switches} &= p \cdot \log_2 p \\ b = p, u = \log_2 p \end{aligned}$$

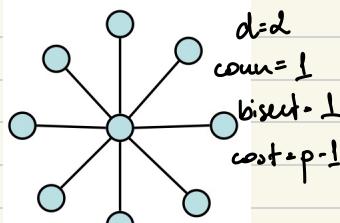
perfect shuffle:  
 $j = \sum_{i=0}^{d-1} i \cdot 2^i, 0 \leq i \leq \frac{p}{2}-1$   
 $\lceil \frac{p}{2} \rceil - p, \frac{p}{2} \leq u \leq p-1$

## Omega Network (dynamic)

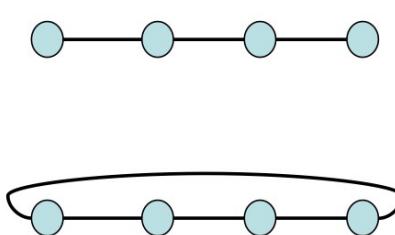


$\beta_1, \beta_2, \dots, \beta_k$   
 $d = \log_2 p, \text{ count} = \frac{p}{2}$   
at stage  $k$ :  $\text{bisection} = \frac{p}{2}$   
 $\beta_k = 0 \rightarrow \text{upper link of switch}$   
 $\beta_k = 1 \rightarrow \text{lower link of switch}$   
processors:  $p$   
stages:  $\log_2 p$   
switches per stage:  $\frac{p}{2}$   
# switches:  $\frac{p}{2} \log_2 p$

## Star Network (static)

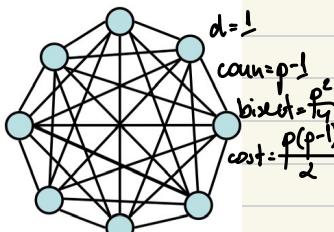


## Link Arrays and Links (static)

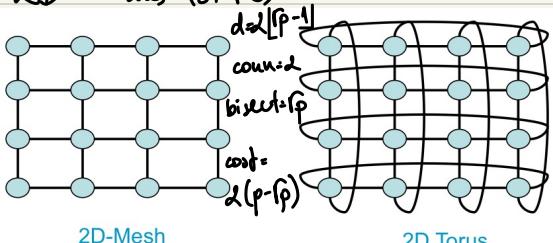


$d = \lfloor \frac{p}{2} \rfloor$   
count = 2  
bisection = 2  
cost =  $p$

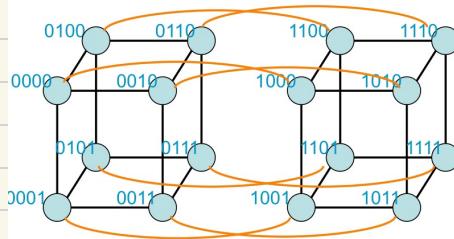
## Completely Connected (static)



## 2D Meshes (static)



## 4-D Hypercube

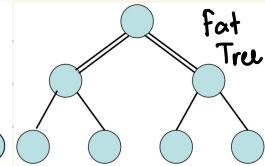
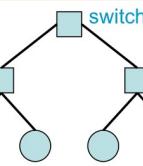
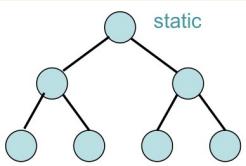


$p = d^d$ ,  $d$ -dimension  
#connections/processor =  $\log p$

max distance =  $\log p$

$d = \log p$ ,  $\text{conn} = \log p$ ,  $\text{bisection} = \frac{p}{2}$   
 $\text{cost} = \frac{p \log p}{2}$

## Tree Network



$d = 2 \log p$   
 $\text{conn} = 1$   
 $\text{bisection} = 1$   
 $\text{cost} = p - 1$

## Evaluation Criteria Static Networks

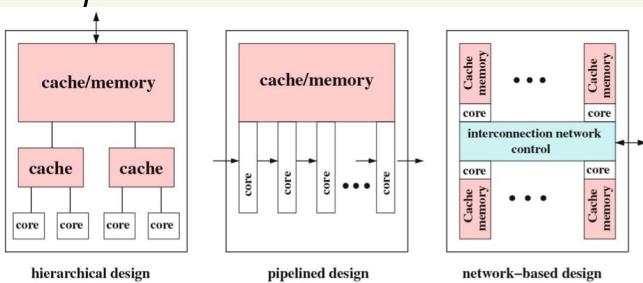
diameter = max distance

connectivity = min # links removed to split in two disjoint parts

bisection width = min # links removed to split in two equal parts

cost = total # links

## Multicore processor



# Lecture 3 - Process, Thread & OpenMP

## Parallelization of programs

- 1) decomposition of computations (partitioning)
- 2) assignment of tasks (scheduling)
- 3) mapping of processes to physical processors

## Data dependencies

- flow (true) dependency - output of  $S_1$  is input of  $S_2$
- anti-dependency - input of  $S_1$  overwritten by  $S_2$
- output dependency -  $S_1, S_2$  store outputs at same place

$$I_1: R_1 \leftarrow R_2 + R_3$$

$$I_2: R_5 \leftarrow R_1 + R_4$$

flow dependency

$$I_1: R_1 \leftarrow R_2 + R_3$$

$$I_2: R_2 \leftarrow R_4 + R_5$$

anti dependency

$$I_1: R_1 \leftarrow R_2 + R_3$$

$$I_2: R_1 \leftarrow R_4 + R_5$$

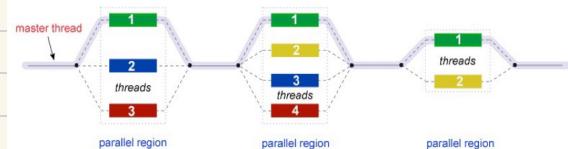
output dependency

Data parallelism - distribute the data structure evenly among processors  
and each perform the operation on its assigned elements  
process - program in execution with own local address space (distributed mem)  
thread - sequence of instructions (shared memory)

## Multi-threading

- timesliced - processor/core switch threads after time interval elapsed
- simultaneous - several threads executed at same time by multiple processors/cores

## OpenMP - programming model for shared memory systems



## for-construct schedule clause:

- parameter chunk - divide iterations into blocks of size chunk
- > static-blocks are assigned to threads in a round-robin fashion (default)
- default chunk = #iterations / #threads or #threads = #blocks
- > dynamic-blocks are assigned to threads in order in which threads finish previous blocks
- default chunk = 1
- > guided-block size starts with implementation-defined value, then decreases exponentially to chunk

## Data Scoping:

- > private - variables local to each thread (new uninitialized instance is created)
- loop control, non-static local to Parallel Region (variables)
- > firstprivate - initialization with Master's value
- > lastprivate - value of last loop iteration is written back to Master
- > shared - variables shared among all threads
- static variables, general default for Parallel Region

Critical Region - executed by all threads, but only by one thread simultaneously  
(Mutual Exclusion)

barrier - wait until all threads have reached the barrier

worksharing constructs - implicit barrier at the end

explicit nowait clause - remove implicit barrier at the end

reduction (operations: list) - +, -, \*, &, |, &&, ||, ^, max, min

## Cache Efficiency

C - row-major storage



a[8][8] : a[0,0], a[0,1], a[0,2], ..., a[0,7], a[1,0], a[1,1], ..., a[1,7], a[2,0], ..., a[2,7]

64B cache-line size => fetch a[0,0] => a[0,0], a[0,1], ..., a[0,7] (8x8B double) fetched

temporal locality - data is used frequently

spatial locality - consecutive data on same cache-line

↳ different threads use elements of same cache-line

False Sharing ↳ one thread writes to cache-line

## Lecture 4 - Performance Analysis

User CPU time of A - time CPU spends executing A

System CPU time of A - time CPU spends executing A + time of operating system

Wall-clock time (waiting time) of A - time elapsed from beginning to end (incl. I/O, execution of others)

Parallel execution time of A - incl. thread synchronization and communication time

$$T_{u\text{-cpu}}(A) = n_{\text{cycle}}(A) \cdot \text{cycle} \rightarrow \text{clock cycle time of CPU}$$

↳ number of CPU cycles for executing A

$$\text{Memory access: } T_{u\text{-cpu}}(A) = (n_{\text{cycle}}(A) + n_{\text{mem\_cycle}}(A)) \cdot \text{cycle}$$

number of additional cycles caused by memory access of A

Cache miss - reading content not residing in cache  $\Rightarrow$  extra delay loading cache line from memory

$$n_{\text{read\_cycles}}(A) = n_{\text{read\_op}}(A) \cdot \text{read\_miss}(A) \cdot n_{\text{miss\_cycle}}$$

# read operations  $\hookrightarrow$  read miss rate  $\hookrightarrow$  # cycles to load cache line

$$T_{u\text{-cpu}}(A) = n_{\text{inst}}(A) \cdot (\text{CPS}(A) + f_{\text{rw\_op}} \cdot r_{\text{miss}}(A) \cdot n_{\text{miss\_cycle}}) \cdot \text{cycle}$$

Clock cycles per Instruction  $\hookrightarrow$  percentage of instructions executing read/write operations

$T_s$  - Serial execution time

$T_p$  - Parallel execution time

$T_o = p \cdot T_p - T_s$  - Overhead function

$S = T_{\text{serial,best}} / T_p$  - True Speed up

$S = T_{p=1} / T_p$  - Relative Speed up

$S_{\text{linear}} = p$  - Linear Speedup

$S_{\text{sublinear}} < p$  - Sublinear Speedup

$S_{\text{superlinear}} > p$  - Superlinear Speedup

Superlinear Speedup  $\downarrow$  FLOP per memory access

Cache: 10ns  $\downarrow t_{\text{access}} = p_{\text{cache}} \cdot 10 + (1-p_{\text{cache}}) \cdot 100$

Main Memory: 100ns

sequential:  $p_{\text{cache}} = 80\% \Rightarrow t_{\text{access}} = 18\text{ ns} \Rightarrow \text{speed} \approx 36 \text{ MFLOP}$

parallel:  $p_{\text{cache}} = 90\% \Rightarrow t_{\text{access}} = 19\text{ ns} \Rightarrow \text{speed per processor} \approx 53 \text{ MFLOP}$

$$S = \frac{2.53}{36} \approx 3$$

Adding  $n$  numbers on  $n$  processors

$$W(n) = n + \sum_{j=1}^{\log n} \lfloor n/2^j \rfloor + 1 \approx 2n + 1 = \Theta(n)$$

$$T(n) = \lg n + d = \Theta(\lg n)$$

Adding  $n$  numbers on  $p$  processors

$$T_p = \frac{n-1}{p} + d \lg p$$

$$\lim_{n \rightarrow \infty} S_p(n) = \lim_{n \rightarrow \infty} \frac{n-1}{p + d \lg p} = p$$

$$E = \frac{S}{p}: E = \frac{1}{1 + \frac{d \lg p}{n-1}}$$

Amdahl's Law

$$W = T_{\text{ser}} + T_{\text{par}} \quad | \quad S = \frac{p}{1 + d(p-1)}, \quad d = \frac{1 - f}{1 - f}$$

**Input:**  $n = 2^k$  numbers stored in array  $A$

**Output:** The sum  $S = \sum_{i=1}^n A[i]$

1. **for**  $1 \leq i \leq n$  **do parallel**

$$B[i] := A[i]$$

2. **for**  $h = 1$  to  $\log n$  **do sequential**

**for**  $1 \leq i \leq n/2^h$  **do parallel**

$$B[i] := B[2i-1] + B[2i]$$

3.  $S := B[1]$

Scalability

$$E = \frac{S}{P} = \frac{T_s}{P \cdot T_p} = \frac{1}{1 + \frac{T_o}{T_s}}$$

$$\text{Gustafsson's law: } \lim_{n \rightarrow \infty} S_p(n) = \lim_{n \rightarrow \infty} \frac{n-1}{p + d \lg p} = p$$

$$\text{Amdahl's law: } \lim_{p \rightarrow \infty} S = \frac{p}{1 + d(p-1)} = f$$

Problem size  $W(n)$  - #basic computation steps in best sequential algorithm

$\xi$ -iso-efficiency

$$T_o = p \cdot T_p - W, \quad S = \frac{W}{T_p} = \frac{W \cdot p}{W + p} \Rightarrow E = \frac{1}{1 + T_o/W}$$

$$W = \left( \frac{E}{1-E} \right) \cdot T_o = K \cdot T_o (W, p) \quad \text{with} \quad K = \frac{E}{1-E}$$

$$W = T_s = n^3 + n \quad | \quad T_o = p T_p - W = 2n p \lg p - n$$

$$T_p = \frac{n^3}{p} + 2n \lg p \quad | \quad T_o/W \approx 2p \lg p / n^3 \Rightarrow \text{const. } E \text{ iff } p = x, \quad n = \sqrt[3]{2x^4/p^3} \times$$

Parallel scalar (dot) product:  $a \cdot b = \sum_{i=1}^n a_i b_i \Rightarrow \frac{1}{p}$  scalar products across  $p$  processors

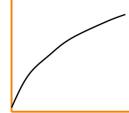
$$T(p, n) = \left( \frac{dn}{p-1} \right) d + \left( \frac{p}{p-1} \right) (d+p) \approx \frac{2n}{p} d + \frac{p}{2} (d+p)$$

$$W = T_s = (2n-1)d \approx dn d \quad | \quad \text{Iso-efficiency: } \frac{T_o}{W} = \frac{p^2 (d+p)}{4nd} = \text{constant}$$

$$T_o = p T(p, n) - T_s = \frac{p^2}{dn} (d+p)$$



If  $p$  goes up,  $E$  goes down



If  $p$  remains constant and the problem size goes up, then  $E$  goes up

## Lecture 5 - MPI

**MPJ** program = set of processes each with own local data

Principle: MPMD (Multiple Program Multiple Data)

Practice: SPMD (Single Program Multiple Data)

Data exchange through messaging over the network (explicit (MPI) or implicit (PGAs))

Process - running in-memory instance of an executable file

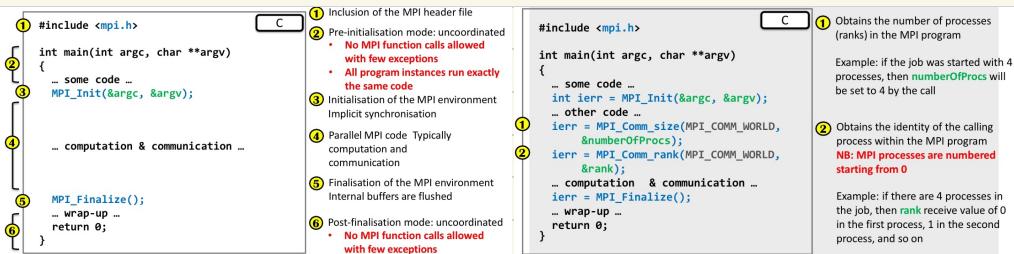
→ executable code

→ one or more threads of execution (sharing memory address space)

→ memory: data, heap, stack, processor state (CPU registers and flags)

→ operating system context

→ ~~PSD~~

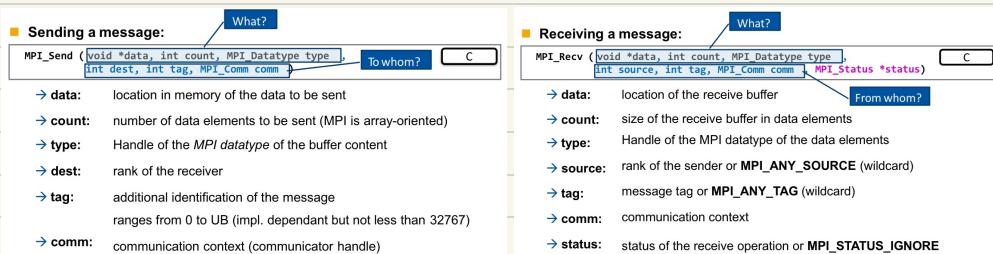


Explicit messaging:

→ send and receive primitives (operations)

→ known addresses of both sender and receiver

→ specification of what has to be sent/received



MPI data type	C data type	
MPI_CHAR	char	send = receive ✓
MPI_SHORT	short	send > receive X
MPI_INT	int	C: MPI_Status status
MPI_FLOAT	float	→ status.MPI_SOURCE message source rank
MPI_DOUBLE	double	→ status.MPI_TAG message tag
MPI_UNSIGNED_INT	unsigned int	→ status.MPI_ERROR receive status code

```
MPI_Sendrecv (void *senddata, int sendcount, MPI_Datatype sendtype,
              int dest, int sendtag, void *recvdata, int recvcount,
              MPI_Datatype recvtype, int source, int recvtag,
              MPI_Comm comm, MPI_Status *status)
```

	Send	Receive
Data	senddata	recvdata
Count	sendcount	recvcount
Type	sendtype	recvtype
Destination	dest	-
Source	-	source
Tag	sendtag	recvtag
Communicator	comm	comm
Receive status	-	status

↳ send 1 message + receive 1 message (no deadlock)  
 ↳ send and receive buffer shouldn't overlap

Non-blocking:

```
MPI_Isend (void *data, int count, MPI_Datatype dataType,
            int dest, int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_Irecv (void *data, int count, MPI_Datatype dataType,
            int source, int tag, MPI_Comm comm, MPI_Request *request)
```

Blocking:

```
MPI_Wait (MPI_Request *request, MPI_Status *status)
```

Test:

```
MPI_Test (MPI_Request *request, int *flag, MPI_Status *status)
```

MPI\_Waitany / MPI\_Testany

- Wait for one of the specified requests to complete and free it
- Test if one of the specified requests has completed and free it if it did

MPI\_Waitall / MPI\_Testall

- Wait for all the specified requests to complete and free them
- Test if all of the specified requests have completed and free them if they have

MPI\_Waitsome / MPI\_Testsome

- Wait for any number of the specified requests to complete and free them
- Test if any number of the specified requests have completed and free these that have

## Send Modes

### Standard mode

→ The call blocks until the message has either been transferred or copied to an internal buffer for later delivery

### Synchronous mode

→ The call blocks until a matching receive has been posted and the message reception has started

### Buffered mode

→ The call blocks until the message has been copied to a user-supplied buffer. Actual transmission may happen at a later point

### Ready mode (don't use!)

→ The operation succeeds only if a matching receive has already been posted. Behaves as standard send in every other aspect

→ MPI\_Send  
→ MPI\_Isend  
→ MPI\_Ssend  
→ MPI\_Issend  
→ MPI\_Bsend  
→ MPI\_Ibsend  
→ MPI\_Rsend  
→ MPI\_Irsend

blocking standard send  
non-blocking standard send  
blocking synchronous send  
non-blocking synchronous send  
blocking buffered send  
non-blocking buffered send  
blocking ready-mode send  
non-blocking ready-mode send

→ MPI\_Buffer\_attach (void \*buf, int size)  
→ MPI\_Buffer\_detach (void \*buf, int \*size)

MPI\_Abort (MPI\_Comm comm, int errorcode)

## double MPI\_Wtime ()

MPI\_Get\_processor\_name (char \*name, int \*resultlen)

## MPI\_Reduce

MPI\_Reduce( void\* send\_data, void\* recv\_data, int count,  
MPI\_Datatype datatype, MPI\_Op op, int root, MPI\_Comm  
communicator)

The reduction operations defined by MPI include:

- MPI\_MAX - Returns the maximum element.
- MPI\_MIN - Returns the minimum element.
- MPI\_SUM - Sums the elements.
- MPI\_PROD - Multiplies all elements.
- MPI\_LAND - Performs a logical *and* across the elements.
- MPI\_LOR - Performs a logical *or* across the elements.
- MPI\_BAND - Performs a bitwise *and* across the bits of the elements.
- MPI\_BOR - Performs a bitwise *or* across the bits of the elements.
- MPI\_MAXLOC - Returns the maximum value and the rank of the process that owns it
- MPI\_MINLOC - Returns the minimum value and the rank of the process that owns it.

# Lecture 6 - Parallel Algorithms for Systems of Linear Equations

## Gaussian Elimination

$$\begin{array}{l} a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n = b_1 \\ \vdots \\ a_{ii}x_i + a_{i2}x_2 + \dots + a_{in}x_n = b_i \\ \vdots \\ a_{nn}x_n + a_{n2}x_2 + \dots + a_{nn}x_n = b_n \end{array}$$

For  $k=1, \dots, n$ , eliminate the  $k$ -th lower column of  $A^{(k)}$

$$l_{ik} = a_{ik}^{(k)} / a_{kk}^{(k)}, \quad i = k+1, \dots, n.$$

Update  $A^{(k)}$ :

$$a_{ij}^{(k+1)} = a_{ij}^{(k)} - l_{ik}a_{kj}^{(k)}$$

$$b_i^{(k+1)} = b_i^{(k)} - l_{ik}b_k^{(k)}$$

Backward substitution:

$$x_n, x_{n-1}, \dots, x_1$$

$$x_k = \frac{1}{a_{kk}^{(n)}} \left( b_k^{(n)} - \sum_{j=k+1}^n a_{kj}^{(n)} x_j \right)$$

LU decomposition,  $A = LU$

$$U = A^{(n)}, L = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}, Ax = b, Ly = b$$

solve  $y$ :  $Ly = b$ , forward substitution

solve  $x$ :  $Ux = y$ , backward substitution

$$A^{(1)} = A, \text{ for } k \geq 2: \quad A^{(k)} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1,k-1} & a_{1k} & \cdots & a_{1n} \\ 0 & a_{22}^{(2)} & \cdots & a_{2,k-1}^{(2)} & a_{2k}^{(2)} & \cdots & a_{2n}^{(2)} \\ \vdots & \ddots & \ddots & \vdots & \vdots & \ddots & \vdots \\ \vdots & \ddots & \ddots & a_{k-1,k-1}^{(k-1)} & a_{k-1,k}^{(k-1)} & \cdots & a_{k-1,n}^{(k-1)} \\ \vdots & & 0 & a_{kk}^{(k)} & \cdots & a_{kn}^{(k)} \\ \vdots & & & \vdots & \ddots & \vdots \\ 0 & \cdots & \cdots & 0 & a_{nk}^{(k)} & \cdots & a_{nn}^{(k)} \end{bmatrix}.$$

## Recursive Doubling

$$A = \begin{pmatrix} a_1 & c_1 & & & 0 \\ a_2 & b_2 & c_2 & & \\ & & & \ddots & \\ a_3 & b_3 & & \ddots & \\ & & \ddots & \ddots & c_{n-1} \\ 0 & & & a_n & b_n \end{pmatrix}$$

Solve  $Ax = y$

Step  $k$ :

1. Compute  $l_{ik} := a_{ik}^{(k)} / a_{kk}^{(k)}$  for  $i = k+1, \dots, n$ .
2. Subtract  $l_{ik}$  times the  $k$ th row from the rows  $i = k+1, \dots, n$ , i.e., compute

$$a_{ij}^{(k+1)} = a_{ij}^{(k)} - l_{ik} \cdot a_{kj}^{(k)} \quad \text{for } k \leq j \leq n \text{ and } k < i \leq n.$$

$$y_i^{(k+1)} = y_i^{(k)} - l_{ik} \cdot y_k^{(k)}$$

Considering the 3 neighboring equations  $i-1, i$ , and  $i+1$  ( $i=3, 4, \dots, n-2$ ):

$$\begin{aligned} a_{i-1}x_{i-2} + b_{i-1}x_{i-1} + c_{i-1}x_i &= y_{i-1} \\ a_{i-1}x_{i-1} + b_i x_i + c_i x_{i+1} &= y_i \\ a_{i+1}x_i + b_{i+1}x_{i+1} + c_{i+1}x_{i+2} &= y_{i+1} \end{aligned}$$

In step 2, consider the equations  $i-2, i, i+2$  of the  $A^{(1)}x = y^{(1)}$  for  $i = 5, 6, \dots, n-2$ . Use equation  $i-2$  to eliminate  $x_{i-2}$  from the  $i$ th equation and use equation  $i+2$  to eliminate  $x_{i+2}$  from the  $i$ th equation.

Use equation  $i-1$  to eliminate  $x_{i-1}$  from the  $i$ th equation and use

equation  $i+1$  to eliminate  $x_{i+1}$  from the  $i$ th equation. (Step 1)

The resulting new equation  $i$  is:

$$a_i^{(1)}x_{i-2} + b_i^{(1)}x_i + c_i^{(1)}x_{i+2} = y_i^{(1)}$$

$$A^{(1)} = \begin{pmatrix} b_1^{(1)} & 0 & c_1^{(1)} & & 0 \\ 0 & b_2^{(1)} & 0 & c_2^{(1)} & \\ & & & & c_2^{(2)} \\ a_3^{(1)} & 0 & b_3^{(1)} & \ddots & \\ & & & \ddots & \ddots \\ a_4^{(1)} & \ddots & \ddots & \ddots & c_{n-2}^{(1)} \\ & \ddots & \ddots & \ddots & 0 \\ 0 & & a_n^{(1)} & 0 & b_n^{(1)} \end{pmatrix}$$

$$A^{(2)} = \begin{pmatrix} b_1^{(2)} & 0 & 0 & 0 & c_1^{(2)} & 0 \\ 0 & b_2^{(2)} & & & c_2^{(2)} & \\ & & \ddots & & \ddots & \\ & & & \ddots & & \\ 0 & & & & & c_{n-4}^{(2)} \\ a_5^{(2)} & & & & & 0 \\ & a_6^{(2)} & & & & 0 \\ & & \ddots & & & 0 \\ 0 & & a_n^{(2)} & 0 & 0 & b_n^{(2)} \end{pmatrix}$$

Step  $k$ : Considering equations  $i-2^{k-1}, i, i+2^{k-1}$ :

Use equation  $i-2^{k-1}$  to eliminate  $x_{i-2^{k-1}}$  from the  $i$ th equation and use equation  $i+2^{k-1}$  is used to eliminate  $x_{i+2^{k-1}}$  from the  $i$ th equation.

## Cyclic Reduction

1. Elimination phase: For  $k = 1, \dots, \lfloor \log n \rfloor$  compute  $a_i^{(k)}, b_i^{(k)}, c_i^{(k)}$  and  $y_i^{(k)}$  with  $i = 2^k, \dots, n$  and step-size  $2^k$ . The number of equations of the form (8.26) is reduced by a factor of 1/2 in each step. In step  $k = \lfloor \log n \rfloor$  there is only one equation left for  $i = 2^N$  with  $N = \lfloor \log n \rfloor$ .

$$a_i^{(k)} x_{i-2^k} + b_i^{(k)} x_i + c_i^{(k)} x_{i+2^k} = y_i^{(k)} \quad (8.26)$$

2. Substitution phase: For  $k = \lfloor \log n \rfloor, \dots, 0$  compute  $x_i$  according to Eq. (8.26) for  $i = 2^k, \dots, n$  with step-size  $2^{k+1}$ :

$$x_i = \frac{y_i^{(k)} - a_i^{(k)} \cdot x_{i-2^k} - c_i^{(k)} \cdot x_{i+2^k}}{b_i^{(k)}}. \quad (8.27)$$

## Lecture 7 - Parallel Iterative Algorithms for Linear Systems

$A = M - N$  s.t.  $M^{-1}$  is easily computed

$$M \xrightarrow{k+1} N x^k + b \quad \text{or} \quad x^{k+1} = C x^k + d, \quad C = M^{-1} N, \quad d = M^{-1} b$$

$A = D - L - R, \quad D = \text{diag}(A)$

$$\text{Jacobi: } D \bar{x}^{k+1} = (L + R) x^k + b, \quad C_{\text{Ja}} = D^{-1} (L + R) = D^{-1} (A - D)$$

$$\text{Gauss-Seidel: } (D - L) \bar{x}^{k+1} = R x^k + b, \quad C_{\text{GS}} = (D - L)^{-1} R$$

$$\text{JOR: } A = \frac{D}{\omega} - L - R - \frac{1-\omega}{\omega} D$$

$$\text{SOR: } (D - \omega L) \bar{x}^{k+1} = (1-\omega) D x^k + \omega R x^k + \omega b$$

$$\begin{aligned} \text{Jacobi: } & A \bar{x} = b \\ \bar{x}^k = D^{-1} \cdot (b - (A - D) \bar{x}^{k-1}) &= (I - D^{-1} \cdot A) \cdot \bar{x}^{k-1} + D^{-1} \cdot b \end{aligned}$$

## Design Issues

→ Optimality: minimal parallel execution time

↳ Trade-off: maximal load balance vs. minimal communication overhead

→ Data Locality: computing time b/w consecutive communications  
communication time (overhead)

→ Grain Size: computing time b/w consecutive communications

## Data Locality Ratios

$A, B \in \mathbb{R}^{n \times n}$ ,  $\vec{x} \in \mathbb{R}^n$ ,  $c$ -time of 1 computation,  $\tau$ -time of communication of 1 number

Matrix-vector product  $A\vec{x}$ :

$$\frac{\frac{n}{p}(2n-1)c}{(\frac{n^2}{p}+n)\tau} \approx \frac{2c}{\tau}$$

Matrix-matrix product  $A \cdot B$ :

$$\frac{\frac{n^2}{p}(2n-1)c}{n^2(\frac{1}{p}+1)\tau} \approx \frac{2nc}{p\tau}$$

(Gauss-Seidel

$$\vec{x}^k = D^{-1}B^* - D^{-1} \cdot (\lambda \vec{x}^k + R \vec{x}^{k-1})$$

D-diagonals, L-lower triangular, R-upper triangular