

Lecture 1 - Introduction to Heuristic Search

agent - entity that perceives its environment and acts upon it

state - a configuration of the agent and its environment

initial state - the state in which the agent begins

actions - choices that can be made in a state

f : Actions(s): return the set of actions that can be executed in state s

transition model - a description of what state results from performing any applicable action in any state

f : Result(s,a): return the state resulting from performing action a in state s

state space - the set of all states reachable from the initial state by any sequence of actions

goal test - way to determine whether a given state is a goal state

path cost - numerical cost associated with a given path

DFS - always explore deepest node in frontier (stack = LIFO)

→ memory efficient

→ stuck in deep solutions

→ worst-case: visit all nodes

BFS - always explore shallowest node in frontier (queue = FIFO)

→ finds shortest path (if unweighted)

→ memory intensive

→ worst-case: visit all nodes

Iterative Deepening

- depth-bounded DFS

Informed Search:

(Greedy) Best-First Search

→ always explore best node in frontier

→ measure by heuristic $h(u)$

→ priority queue

Algorithm	Complete? finds soln if one exists	Time	Space	Optimal? finds shortest path
DFS	if $D < \infty$	$O(b^D)$	$O(b \cdot D)$	-
BFS	✓	$O(b^d)$	$O(b^d)$	if unweighted
IterDFS	✓	exercise	exercise	-
BestFS	if $D < \infty$	$O(b^D)$	$O(b \cdot D)$	-
Dijkstra	✓	$O(b^e)$	$O(b^e)$	✓
A*	if $D < \infty$	$O(b^e)$	$O(b^e)$	✓

b : (max) branching factor of tree D : depth of tree d : depth of solution
 e : effective depth, depends on heuristic

Lecture 2 - A* Search

Informed Search - heuristic function estimates forward cost
Dijkstra = Uniform Cost Search

→ cost-from-start-to-current-node → guaranteed shortest path

A* Search - look back (path cost) and forward (heuristic estimate)

→ Disadvantage: entire explored region in memory

Heuristic Properties

→ admissible - never over-estimate the true cost

→ consistent - never over-estimate the growth of the path cost

admissible: $0 \leq h(n) \leq h^*(n)$, consistent: $h(n) \leq h(n') + c(n, n')$, consistent → admissible

Optimality - finds the shortest path

→ Tree Search: admissible heuristic → A* - optimal, $h=0 \rightarrow A^* \equiv UCS$

→ Graph Search: consistent → A* optimal, UCS - optimal

Optimally efficient - does least work for any such algorithm

consistent heuristic → A* - optimally efficient

Efficiency - set of nodes expanded (not # expansions)

Dominance: $\forall n: h_1(n) \geq h_2(n) \rightarrow h_1 \geq h_2$

→ larger is better (if still admissible)

→ trivial heuristic is worst (admissible)

→ exact heuristic is best, but expensive

Non-dominating heuristics: $h'(n) = \max(h_1(n), h_2(n))$ - admissible and dominant

Algorithm	Complete? <small>finds soln if one exists</small>	Time	Space	Optimal? <small>finds shortest path</small>
DFS	if $D < \infty$	$O(b^D)$	$O(b \cdot D)$	—
BFS	✓	$O(b^d)$	$O(b^d)$	if unweighted
IDDFS	✓	exercise	exercise	—
BestFS	if $D < \infty$	$O(b^D)$	$O(b \cdot D)$	—
Dijkstra	✓	$O(b^e)$	$O(b^e)$	✓
A*	if $D < \infty$	$O(b^e)$	$O(b^e)$	if h admissible
IDA*	if $D < \infty$	$O(b^e)$	$O(b \cdot D)$	if h admissible

b : (max) branching factor of tree D : depth of tree d : depth of solution
 e : effective depth, depends on heuristic

Lecture 3 - Adversarial Search

Terminology of Games

single / two / multi-player
Solitaire / Chess / Catan

simultaneous/sequential moves
Rock-paper-scissors/Go

stochastic/deterministic moves
Poker / Chess


partial/perfect information
Card-games/connect-four

discrete/continual time
turn-based/real-time strategy

zero-sum / general-sum
Chess / Catan

Minimax Search

Agent plays value-maximizing moves, assuming adversary plays value-minimizing ones.



Improvements:

→ depth-limited

→ pruning

→ stochastic search

evaluation function:

→ efficient neural network (NNUE)

Alpha-Beta Pruning

```

fun minimax(n: node, d: int, min: int, max: int): int =
  if leaf(n) or d=0 return evaluate(n)
  if n is a max node
    v := min
    for each child of n
      v' := minimax (child,d-1,v,max)
      if v' > v, v:= v'
    if v > max return max
  return v
  if n is a min node
    v := max
    for each child of n
      v' := minimax (child,d-1,min,v)
      if v' < v, v:= v'
    if v < min return min
  return v

```

Lecture 4 - Monte Carlo Search Tree

Markov Decision Process

Deterministic episodic two-player MDP $\langle S, A, P, T, R \rangle$

- $\rightarrow s \in S$ - all game-states [discrete, perfect information]
- $\rightarrow a \in A(s)$ - all actions that can be played in state s [discrete, sequential]
- $\rightarrow P(s, a) = s'$ - determines next state $s' \in S$ [deterministic moves]
- $\rightarrow T(s) \in \{ \perp, T \}$ - whether game terminates in s [episodic]
- $\rightarrow R(s) \in \mathbb{R}$ - reward for terminal state s
- $\rightarrow p(s) \in \{ -1, 1 \}$ - which player's move it is [two sequential players]
- $\circ R(s) = +1$ for winning, 0 for draw, -1 for loss

Zero-sum - "one player's gain is the other's loss"

- every reward r for one player is a punishment $-r$ for the other
- two-player zero sum game can be expressed by \pm function
- player 1 maximizes function, player -1 minimizes function

Random Search

Approximate $V(P(s, a))$, $\forall a \in A(s)$ with random rollouts

$$\tilde{V}_R(s) = \frac{1}{N(s)} \sum_{t=1}^{N(s)} r_t, \quad \forall s' \in \{ P(s, a) \mid a \in A(s) \}$$

minimax selection: $a_t := \arg \max_{a \in A(s_t)} p(s_t) \tilde{V}_R(P(s_t, a))$

$$\lim_{N(s) \rightarrow \infty} \tilde{V}_R(s) = \begin{cases} R_s & \text{id } T(s) \\ \frac{1}{|A(s)|} \sum_{a \in A(s)} \tilde{V}_R(P(s, a)) & \text{otherwise} \end{cases}$$

$$UCT = MCTS + UCB$$

exploitation = minimax-like

$$\arg \max_{s'} \left(\frac{w_i}{n_i} + c \sqrt{\ln N_i / n_i} \right)$$

w_i - # wins for s' after i^{th} move

n_i - # simulations for s' after i^{th} move

N_i - # simulations for s after i^{th} move (parent of s')

c - exploration parameter, $\sqrt{2} \rightarrow$ minimax

Monte Carlo Tree Search

partial tree $B \subseteq S$

- \rightarrow minimax-selection in partial tree
- \rightarrow expand tree with new leaf
- \rightarrow approximate lead values with random rollouts
- \rightarrow backup resulting reward to all parents

$$\mu(a|s) - \text{fraction of action } a \text{ chosen in state } s \quad \tilde{V}_\mu(s) = \begin{cases} R(s) & \text{id } T(s) \\ \sum_{a \in A(s)} \mu(a|s) \tilde{V}_\mu(P(s, a)) & \text{otherwise} \end{cases}$$